

Muhang Tian

✉ mt4193@nyu.edu | 📍 60 5th Ave, New York, NY 10011 | 🌐 <https://muhang-tian.com/>

EDUCATION	Courant Institute, New York University Ph.D. in Computer Science	01/2025 – Expected 12/2029
	Duke University B.S. in Computer Science, double minor in Mathematics and Economics <i>Graduation with Highest Distinction, Alex Vasilos Memorial Award</i>	08/2020 – 05/2024

PUBLICATIONS	[1] Reward-rate Policy Gradient for Efficient Machine Learning Engineering Agents	2026
	<u>Muhang Tian</u> , Sherry Yang <i>Preprint</i>	
	[2] Multi-objective Reinforcement Learning with Nonlinear Preferences: Provable Approximation for Maximizing Expected Scalarized Return	2025
	Nianli Peng*, <u>Muhang Tian</u> *, Brandon Fain <i>Proceedings of the International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS)</i>	[Link to paper]
	[3] Reliable Generation of Privacy-preserving Synthetic Electronic Health Record Time Series via Diffusion Models	2024
	<u>Muhang Tian</u> , Bernie Chen, Allan Guo, Shiyi Jiang, Anru R Zhang <i>Journal of the American Medical Informatics Association (JAMIA)</i>	[Link to paper]
	[4] Fast and Interpretable Mortality Risk Scores for Critical Care Patients	2024
	Chloe Qinyu Zhu*, <u>Muhang Tian</u> *, Lesia Semenova, Jiachang Liu, Jack Xu, Joseph Scarpa, Cynthia Rudin <i>Journal of the American Medical Informatics Association (JAMIA)</i>	[Link to paper]
	[5] Welfare and Fairness in Multi-objective Reinforcement Learning	2023
	Zimeng Fan*, Nianli Peng*, <u>Muhang Tian</u> *, Brandon Fain <i>Proceedings of the International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS)</i>	[Link to paper]

WORK EXPERIENCE	ByteDance <i>Research Intern</i>	06/2026 – 09/2026
	Explored an agentic framework to automate research experiments for generative recommender systems.	
	Big Ideas Lab <i>Machine Learning Engineer</i>	05/2022 – 03/2023
	Applied EfficientDet for object detection tasks on COVID-19 test results.	
	Afloat <i>Software Engineer</i>	05/2022 – 08/2022
	Maintained company server, performed software testing, and reported weekly directly to co-founders.	
	Sinolink Securities <i>Equity Research Analyst</i>	05/2021 – 08/2021
	Analyzed market data for target firms, wrote weekly reports, and managed the company's database.	

RESEARCH
EXPERIENCE

K-space Transformer	03/2025 – 03/2026
■ Explored using Vision Transformers to learn meaningful representations from k-space, aiming to perform prediction directly from Fourier transforms of natural images. ■ Explored on ImageNet with different training recipes and normalization schemes. Observed that using a linear layer for tokenization is crucial for accurate prediction.	
Diffusion Model for Synthetic Health Records	05/2023 – 12/2023
■ Explored generating mixed-type time series data with multinomial and Gaussian diffusions using bidirectional recurrent neural network as the backbone. ■ Experimentally demonstrated the proposed approach produces synthetic samples with high data utility, outperforms generative adversarial networks, and improves ML models' performance.	
Nonlinear Multi-objective Reinforcement Learning	05/2022 – 05/2024
■ Explored optimizing the expectation over a nonlinear objective on the cumulative reward. Proposed an approximate algorithm that optimizes this objective by considering past accumulated rewards. ■ Experimentally and theoretically verified that the algorithm converges. Designed another dynamic programming approach to learn a policy with theoretical guarantees.	
Interpretable Machine Learning for Risk Prediction	01/2023 – 10/2023
■ Explored finding an interpretable model for classification task through approximately optimizing a L0 regularized logistic regression model with integer, sparsity, and box constraints. ■ Designed group sparsity and monotonicity constraints to improve readability and usefulness of the risk score in practical settings.	
Bandit and Reinforcement Learning Algorithms	10/2022 – 03/2023
■ Researched Sample Average Uncertainty approach towards multi-arm bandit problems by utilizing approximate sample mean variance as an exploration bonus. ■ Implemented Deep Q-Learning, Proximal Policy Optimization (PPO), and Advantage Actor-Critic algorithms from scratch and compared their performances in Atari environments.	

HONORS AND
AWARDS

Henry M. MacCracken Award: multi-year funding for New York University PhD students.	2024
CS+ Summer Research Fellowship: Duke summer research funding (\$10,000).	2022 – 2023

ACADEMIC
SERVICES

Reviewer: NeurIPS (2023), ICLR (2024), AAAI (2025), IEEE ICHI (2023).
Teaching Assistant: Fundamentals of Machine Learning, Design and Analysis of Algorithms, Econometrics.

TECHNICAL
SKILLS

Languages: Python, Java, JavaScript, C/C++, R, SQL, MATLAB, CSS/HTML.
Softwares: PyTorch, TensorFlow, Keras, scikit-learn, NumPy, Pandas, Linux, Slurm, Docker, Git, Linux.